R Programming Cheat Sheet

JUST THE BASICS

CREATED BY: ARIANNE COLTON AND SEAN CHEN

GENERAL

- R version 3.0 and greater adds support for 64 bit integers
- R is case sensitive
- R index starts from 1

HELP

`help(functionName)` or `?functionName`

Help Home Page	`help.start()`
Special Character Help	`help('[')`
Search Help	`help.search(..)` or `??..`
Search Function - with Partial Name	`apropos('mea')`
See Example(s)	`example(topic)`

OBJECTS in current environment

Display Object Name	`objects()` or `ls()`
Remove Object	`rm(object1, object2,..)`

Notes:
1. .name starting with a period are accessible but invisible, so they will not be found by 'ls'
2. To guarantee memory removal, use 'gc', releasing unused memory to the OS. R performs automatic 'gc' periodically

SYMBOL NAME ENVIRONMENT

- If multiple packages use the same function name the function that the package loaded the last will get called.
- To avoid this precede the function with the name of the package. e.g. `packageName::functionName(..)`

LIBRARY

Only trust reliable R packages i.e., 'ggplot2' for plotting, 'sp' for dealing spatial data, 'reshape2', 'survival', etc.

Load Package	`library(packageName)` or `require(packageName)`
Unload Package	`detach(packageName)`

Note: require() returns the status(True/False)

MANIPULATING STRINGS

Putting Together Strings	`paste('string1', 'string2', sep = '/')`
	 # separator ('sep') is a space by default
	 `paste(c('1', '2'), collapse = '/')`
	 # returns '1/2'

Split String	`stringr::str_split(string = v1, pattern = '-')` # returns a list
Get Substring	`stringr::str_sub(string = v1, start = 1, end = 3)`

Match String	`isJohnFound <- stringr::str_detect(string = df1$col1, pattern = ignore.case('John'))`
	 # returns True/False if John was found
	 `df1[isJohnFound, c('col1', ...)]`

DATA TYPES

Check data type: `class(variable)`

FOUR BASIC DATA TYPES

1. **Numeric** - includes float/double, int, etc.

   `is.numeric(variable)`

2. **Character(string)**

   `nchar(variable)` # length of a character or numeric

3. **Date/POSIXct**
   - **Date**: stores just a date. In numeric form, number of days since 1/1/1970 (see below).

     `date1 <- as.Date('2012-06-28'), as.numeric(date1)`

   - **POSIXct**: stores a date and time. In numeric form, number of seconds since 1/1/1970.

     `date2 <- as.POSIXct('2012-06-28 18:00')`

   Note: Use 'lubridate' and 'chron' packages to work with Dates

4. **Logical**
   - (TRUE = 1, FALSE = 0)
   - Use ==/!= to test equality and inequality

     `as.numeric(TRUE) => 1`

DATA STRUCTURES

VECTOR

- Group of elements of the SAME type
- R is a vectorized language, operations are applied to each element of the vector automatically
- R has no concept of column vectors or row vectors
- Special vectors: letters and LETTERS, that contain lower-case and upper-case letters

Create Vector	`v1 <- c(1, 2, 3)`
Get Length	`length(v1)`
Check if All or Any is True	`all(v1); any(v1)`
Integer Indexing	`v1[1:3]; v1[c(1,6)]`
Boolean Indexing	`v1[is.na(v1)] <- 0`
Naming	`c(first = 'a', ..)` or `names(v1) <- c('first', ..)`

FACTOR

- `as.factor(v1)` gets you the levels which is the number of unique values
- Factors can reduce the size of a variable because they only store unique values, but could be buggy if not used properly

LIST

Store any number of items of ANY type

Create List	`list1 <- list(first = 'a', ...)`
Create Empty List	`vector(mode = 'list', length = 3)`
Get Element	`list1[[1]]` or `list1[['first']]`
Append Using Numeric Index	`list1[[6]] <- 2`
Append Using Name	`list1[['newElement']] <- 2`

Note: repeatedly appending to list, vector, data.frame etc. is expensive, it is best to create a list of a certain size, then fill it.

DATA.FRAME

- Each column is a variable, each row is an observation
- Internally, each column is a vector
- idata.frame is a data structure that creates a reference to a data.frame, therefore, no copying is performed

Create Data Frame	`df1 <- data.frame(col1 = v1, col2 = v2, v3)`
Dimension	`nrow(df1); ncol(df1); dim(df1)`
Get/Set Column Names	`names(df1)` `names(df1) <- c(...)`
Get/Set Row Names	`rownames(df1)` `rownames(df1) <- c(...)`
Preview	`head(df1, n = 10); tail(...)`
Get Data Type	`class(df1)` # is data.frame
Index by Column(s)	`df1['col1']` or `df1[1];` [†] `df1[c('col1', 'col3')]` or `df1[c(1, 3)]`
Index by Rows and Columns	`df1[c(1, 3), 2:3]` # returns data from row 1 & 3, columns 2 to 3

[†] Index method: `df1$col1` or `df1[, 'col1']` or `df1[, 1]` returns as a vector. To return single column

DATA.TABLE

data.frame while using single-square brackets, use 'drop': `df1[, 'col1', drop = FALSE]`

What is a data.table
- Extends and enhances the functionality of data.frames

Differences: data.table vs. data.frame
- By default data.frame turns character data into factors, while data.table does not
- When you print data.frame data, all data prints to the console, with a data.table, it intelligently prints the first and last five rows
- **Key Difference:** Data.tables are fast because they have an index like a database.

  i.e., this search, `dt1$col1 > number`, does a sequential scan (vector scan). After you create a key for this, it will be much faster via binary search.

Create data.table from data.frame	`data.table(df1)`
Index by Column(s)[*]	`dt1[, 'col1', with = FALSE]` or `dt1[, list(col1)]`
Show info for each data.table in memory (i.e., size, ...)	`tables()`
Show Keys in data.table	`key(dt1)`
Create index for col1 and reorder data according to col1	`setkey(dt1, col1)`
Use Key to Select Data	`dt1[c('col1Value1', 'col1Value2'), ]`
Multiple Key Select	`dt1[J('1', c('2', '3')), ]`
Aggregation[**]	`dt1[, list(col1 = mean(col1)), by = col2]` `dt1[, list(col1 = mean(col1), col2Sum = sum(col2)), by = list(col3, col4)]`

[*] Accessing columns must be done via list of actual names, not as characters. If column names are characters, then "with" argument should be set to FALSE.

[**] Aggregate and d*ply functions will work, but built-in aggregation functionality of data table is faster

MATRIX

- Similar to data.frame except every element must be the SAME type, most commonly all numerics
- Functions that work with data.frame should work with matrix as well

Create Matrix	`matrix1 <- matrix(1:10, nrow = 5),` # fills rows 1 to 5, column 1 with 1:5, and column 2 with 6:10
Matrix Multiplication	`matrix1 %*% t(matrix2)` # where t() is transpose

ARRAY

- Multidimensional vector of the SAME type
- `array1 <- array(1:12, dim = c(2, 3, 2))`
- Using arrays is not recommended
- Matrices are restricted to two dimensions while array can have any dimension

DATA MUNGING

APPLY (apply, tapply, lapply, mapply)

- Apply - most restrictive. Must be used on a matrix, all elements must be the same type
- If used on some other object, such as a data.frame, it will be converted to a matrix first

```
apply(matrix1, 1 - rows or 2 - columns,
function to apply)
# if rows, then pass each row as input to the function
```

- By default, computation on NA (missing data) always returns NA, so if a matrix contains NAs, you can ignore them (use `na.rm = TRUE` in the `apply(..)` which doesn't pass NAs to your function)

lapply

Applies a function to each element of a list and returns the results as a list

sapply

Same as lapply except return the results as a vector

> **Note:** lapply & sapply can both take a vector as input, a vector is technically a form of list

AGGREGATE (SQL GROUPBY)

- **aggregate(formulas, data, function)**
- Formulas: y ~ x, y represents a variable that we want to make a calculation on, x represents one or more variables we want to group the calculation by
- Can only use one function in aggregate(). To apply more than one function, use the plyr() package

In the example below diamonds is a data.frame; price, cut, color etc. are columns of diamonds.

```
aggregate(price ~ cut, diamonds, mean)
# get the average price of different cuts for the diamonds
aggregate(price ~ cut + color, diamonds,
mean)  # group by cut and color
aggregate(cbind(price, carat) ~ cut,
diamonds, mean) # get the average price and average
carat of different cuts
```

PLYR ('split-apply-combine')

- ddply(), llply(), ldply(), etc. (1st letter = the type of input, 2nd = the type of output)
- plyr can be slow, most of the functionality in plyr can be accomplished using base function or other packages, but plyr is easier to use

ddply

Takes a data.frame, splits it according to some variable(s), performs a desired action on it and returns a data.frame

llply

- Can use this instead of lapply
- For sapply, can use laply ('a' is array/vector/matrix), however, laply result does not include the names.

DPLYR (for data.frame ONLY)

- Basic functions: filter(), slice(), arrange(), select(), rename(), distinct(), mutate(), summarise(),

group_by(), sample_n()

- Chain functions

```
df1 %>% group_by(year, month) %>%
select(col1, col2) %>% summarise(col1mean
= mean(col1))
```

- Much faster than plyr, with four types of easy-to-use joins (inner, left, semi, anti)
- Abstracts the way data is stored so you can work with data frames, data tables, and remote databases with the same set of functions

HELPER FUNCTIONS

each() - supply multiple functions to a function like aggregate

```
aggregate(price ~ cut, diamonds, each(mean,
median))
```

DATA

LOAD DATA FROM CSV

- Read csv

```
read.table(file = url or filepath, header =
TRUE, sep = ',')
```

- "stringAsFactors" argument defaults to TRUE, set it to FALSE to prevent converting columns to factors. This saves computation time and maintains character data
- Other useful arguments are "quote" and "colClasses", specifying the character used for enclosing cells and the data type for each column.
- If cell separator has been used inside a cell, then use `read.csv2()` or read `delim2()` instead of `read.table()`

DATABASE

Connect to Database	`db1 <- RODBC::odbcConnect('conStr')`
Query Database	`df1 <- RODBC::sqlQuery(db1, 'SELECT ..', stringAsFactors = FALSE)`
Close Connection	`RODBC::odbcClose(db1)`

- Only one connection may be open at a time. The connection automatically closes if R closes or another connection is opened.
- If table name has space, use [ ] to surround the table name in the SQL string.
- which() in R is similar to 'where' in SQL

INCLUDED DATA

R and some packages come with data included.

| List Available Datasets | `data()` |
| List Available Datasets in a Specific Package | `data(package = 'ggplot2')` |

MISSING DATA (NA and NULL)

NULL is not missing, it's nothingness. NULL is atomical and cannot exist within a vector. If used inside a vector, it simply disappears.

| Check Missing Data | `is.na()` |
| Avoid Using | `is.null()` |

FUNCTIONS AND CONTROLS

| Create Function | `say_hello <- function(first, last = 'hola') { }` |
| Call Function | `say_hello(first = 'hello')` |

- R automatically returns the value of the last line of code in a function. This is bad practice. Use return() explicitly instead.
- do.call() - specify the name of a function either as string (i.e. 'mean') or as object (i.e. mean) and provide arguments as a list.

```
do.call(mean, args = list(first = '1st'))
```

IF /ELSE /ELSE IF /SWITCH

	if { } else	ifelse
Works with Vectorized Argument	No	Yes
Most Efficient for Non-Vectorized Argument	Yes	No
Works with NA *	No	Yes
Use &&, || **†	Yes	No
Use &, | ***†	No	Yes

- * NA == 1 result is NA, thus if won't work, it'll be an error. For ifelse, NA will return instead
- ** &&, \|\| is best used in if, since it only compares the first element of vector from each side
- *** &, \| is necessary for ifelse, as it compares every element of vector from each side
- † &&, \|\| are similar to if in that they don't work with vectors, where ifelse, &, \| work with vectors
- Similar to C++/Java, for &, \|, both sides of operator are always checked. For &&, \|\|, if left side fails, no need to check the right side.
- } else, else must be on the same line as }

GRAPHICS

DEFAULT BASIC GRAPHIC

```
hist(df1$col1, main = 'title', xlab = 'x
axis label')
```

```
plot(col2 ~ col1, data = df1),
aka y ~ x or plot(x, y)
```

LATTICE AND GGPLOT2 (more popular)

- Initialize the object and add layers (points, lines, histograms) using +, map variable in the data to an axis or aesthetic using 'aes'

```
ggplot(data = df1) + geom_histogram(aes(x
= col1))
```

- Normalized histogram (pdf, not relative frequency histogram)

```
ggplot(data = df1) + geom_density(aes(x =
col1), fill = 'grey50')
```

DATA RESHAPING

REARRANGE

| Melt Data - from column to row | `reshape2.melt(df1, id.vars = c('col1', 'col2'), variable.name = 'newCol1', value.name = 'newCol2')` |
| Cast Data - from row to column | `reshape2.dcast(df1, col1 + col2 ~ newCol1, value.var = 'newCol2')` |

If df1 has 3 more columns, col3 to col5, 'melting' creates a new df that has 3 rows for each combination of col1 and col2, with the values coming from the respective col3 to col5.

COMBINE (mutiple sets into one)

1. **cbind** - bind by columns

   | data.frame from two vectors | `cbind(v1, v2)` |
   | data.frame combining df1 and df2 columns | `cbind(df1, df2)` |

2. **rbind** - similar to cbind but for rows, you can assign new column names to vectors in cbind

   ```
   cbind(col1 = v1, ...)
   ```

3. **Joins** - (merge, join, data.table) using common keys

   ### 3.1 Merge
   - by.x and by.y specify the key columns use in the join() operation
   - Merge can be much slower than the alternatives

   ```
   merge(x = df1, y = df2, by.x = c('col1',
   'col3'), by.y = c('col3', 'col6'))
   ```

   ### 3.2 Join
   - Join in plyr() package works similar to merge but much faster, drawback is key columns in each table must have the same name
   - join() has an argument for specifying left, right, inner joins

   ```
   join(x = df1, y = df2, by = c('col1',
   'col3'))
   ```

   ### 3.3 data.table

   ```
   dt1 <- data.table(df1, key = c('1',
   '2')),dt2 <- ...‡
   ```

   - Left Join

   ```
   dt1[dt2]
   ```

‡ Data table join requires specifying the keys for the data tables